

Neural Language Processing with PyTorch

Week 1 딥러닝을 위한 PyTorch 실무환경 구축

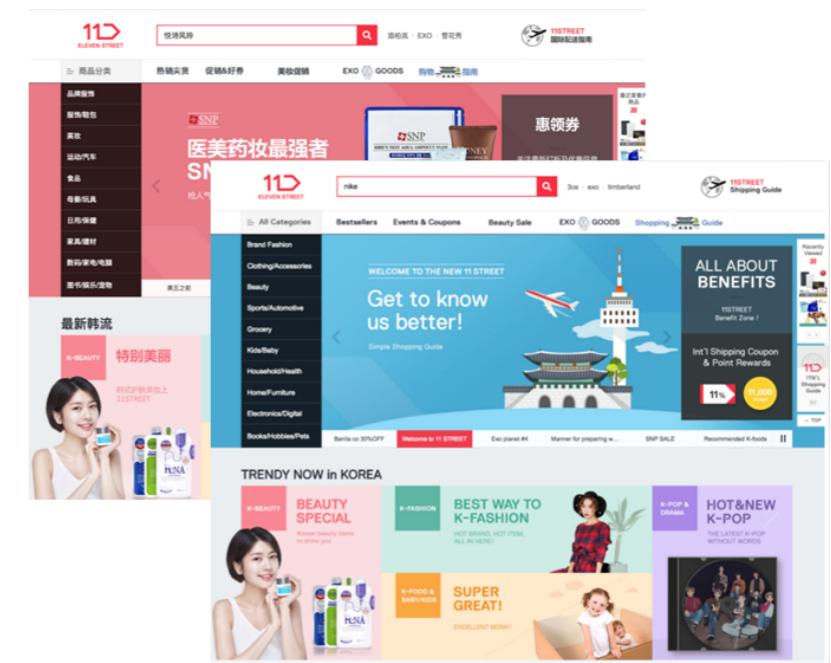
Ki Hyun Kim

- Machine Learning Researcher @ MakinaRocks
- Linkedin: <https://www.linkedin.com/in/ki-hyun-kim/>
- Github: <https://github.com/kh-kim/>
- Email: pointzz.ki@gmail.com



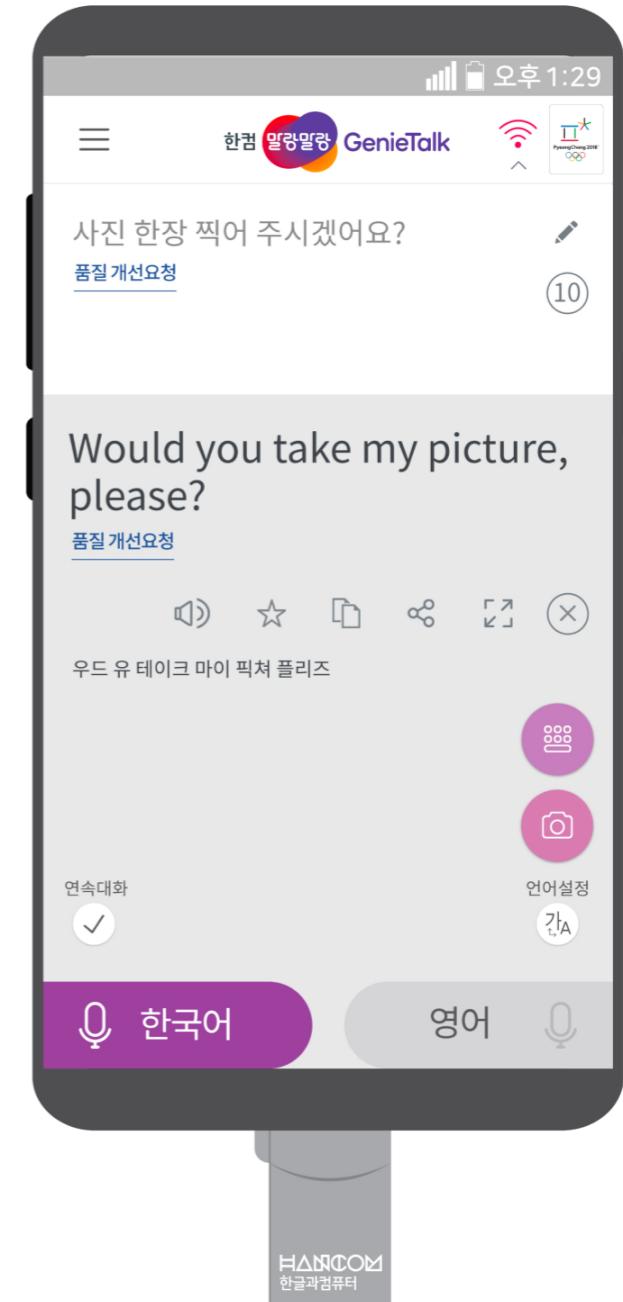
Ki Hyun Kim

- Machine Learning Researcher @ SKPlanet
 - Neural Machine Translation
 - 글로벌 11번가
 - 한영/영한, 한중/중한 기계번역
 - 7000만 개 이상의 상품타이틀 번역, 리뷰 실시간 번역
 - SK AI asset 공유
 - SK C&C Aibril: 한중/중한, 한영/영한, 영중/중영 API 제공
 - SK 그룹 한영중 통번역기 API 제공



Ki Hyun Kim

- Machine Learning Engineer @ TMON
 - Recommender System
 - QA-bot
- Researcher @ ETRI
 - Automatic Speech Translation
 - GenieTalk
- BS + MS of CS @ Stony Brook Univ.



오상준

- Deep Learning Engineer @ Deep Bio
 - 병리영상 기반 전립선암 진단모델 연구개발
 - GPU 서버 분산 스케줄링 시스템 개발
- Co-founder, Research Engineer @ QuantumSurf
 - 선물거래 알고리즘을 위한 API 설계 및 UX 개발
 - IPTV 영상품질 예측모델 연구개발
- BS in English Literature, minor in Philosophy @ 한국외국어대학교



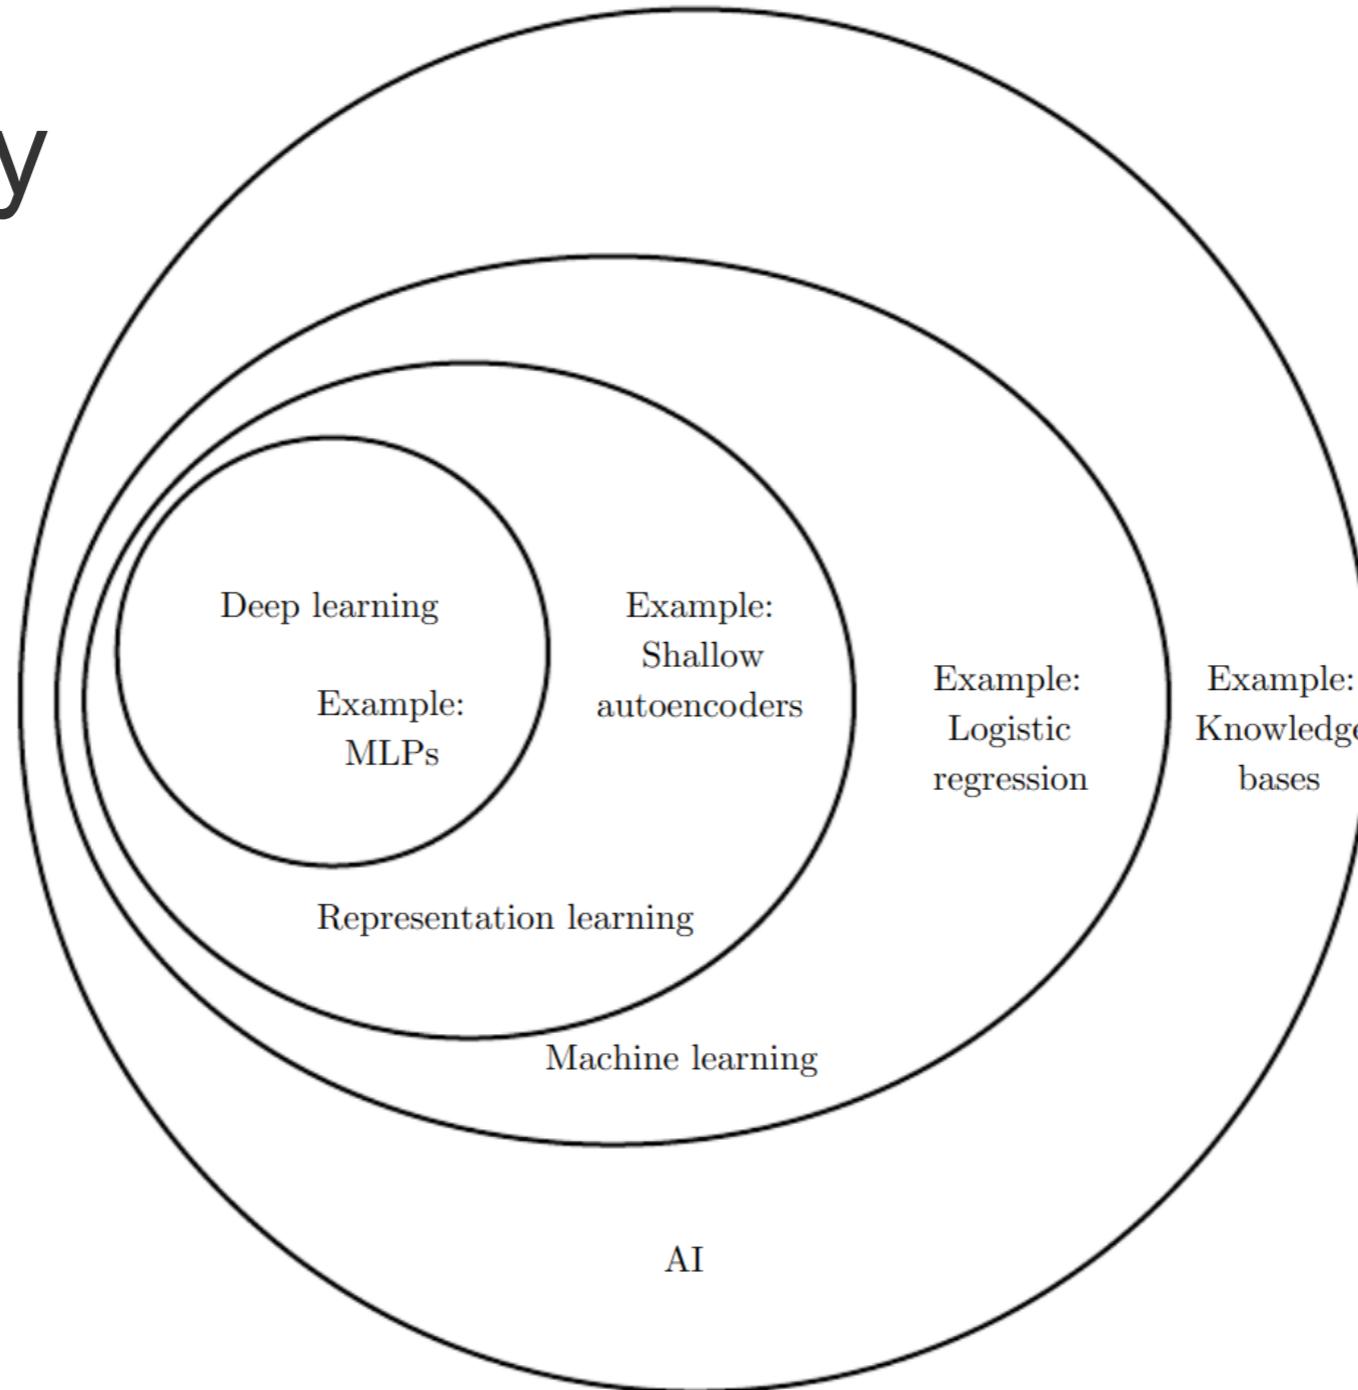
오상준

- Github: <https://github.com/juneoh>
- Email: me@juneoh.net

1. Introduction to Deep Learning

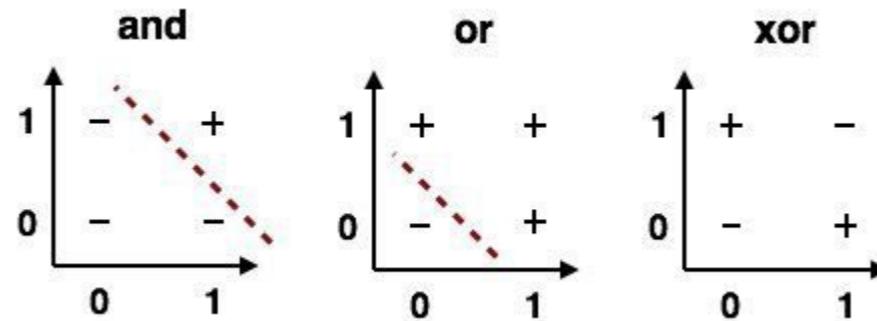


Genealogy



Timeline

- **Cybernetics** 1940s-1960s
 - McCulloh-Pitts neuron
 - McCulloch and Pitts, 1942. A Logical Calculus of the Ideas Immanent in Nervous Activity.
 - Hebbian learning
 - Hebb, 1949. The Organization of Behaviour.
 - Perceptron
 - Rosenblatt, 1958. The Perceptron: A Probabilistic Model for Information Storage and Organization in the Brain.



BRACE YOURSELVES



Timeline

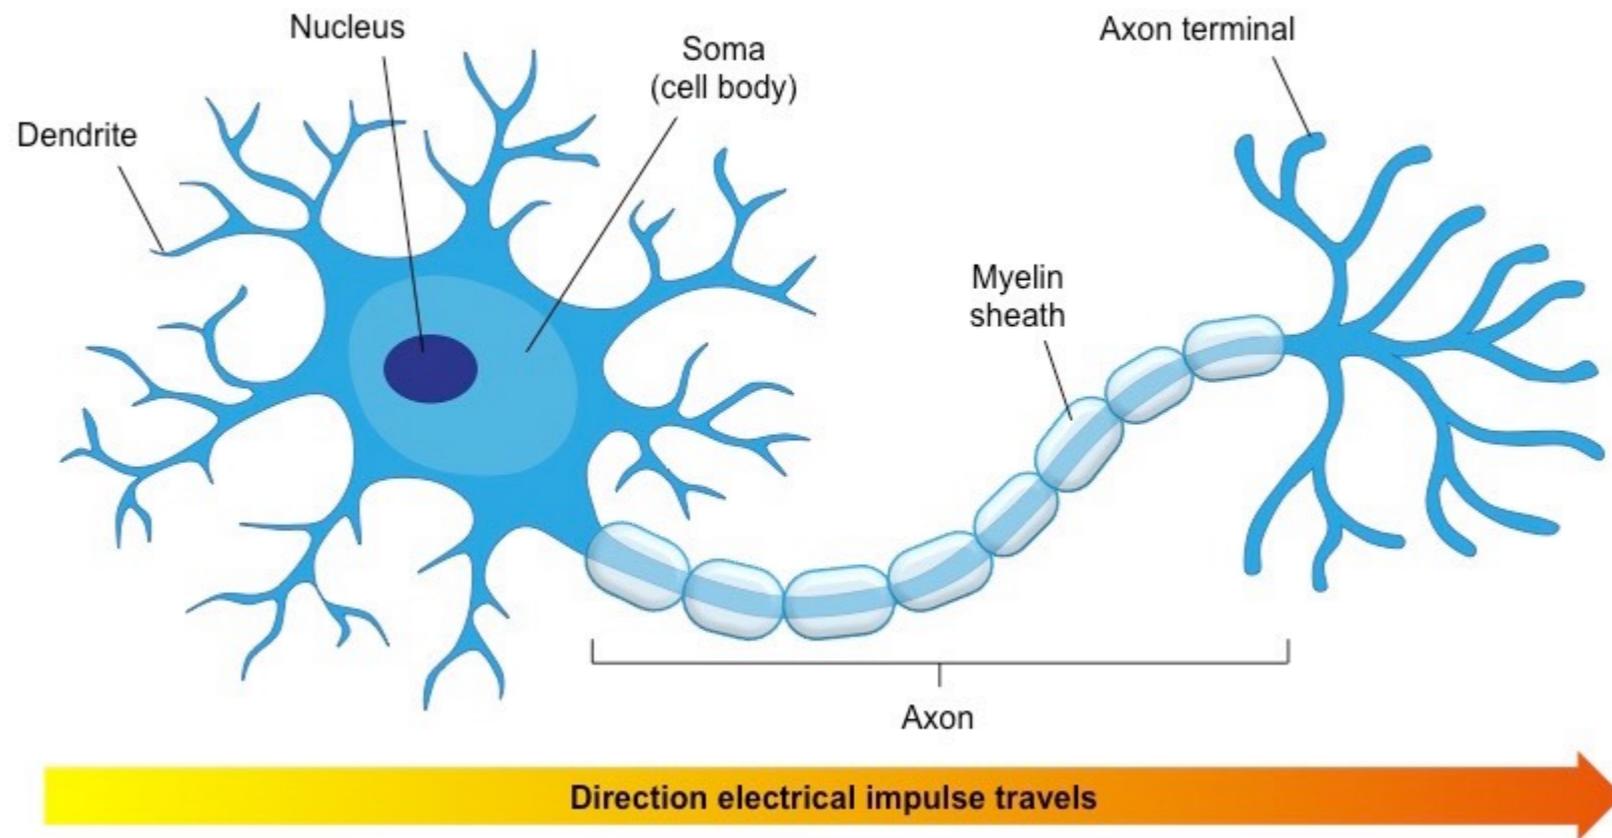
- **Connectionism 1980s-1990s**
 - Backpropagation
 - Rumelhart et al, 1986. Learning Representations by Back-propagating Errors.
 - Convolutional Neural Networks
 - Fukushima, 1980. Neocognitron: A Self-organizing Neural Network Model for a Mechanism of Pattern Recognition Unaffected by Shift in Position.

Timeline

- **Deep Learning 2006-**
 - Deep Neural Networks
 - Hinton et al, 2006. A Fast Learning Algorithm for Deep Belief Nets.
 - Rectified Linear Units
 - Golorot et al, 2011. Deep Sparse Rectifier Neural Networks.
 - AlexNet
 - Krizhevsky et al, 2012. ImageNet Classification with Deep Convolutional Neural Networks.

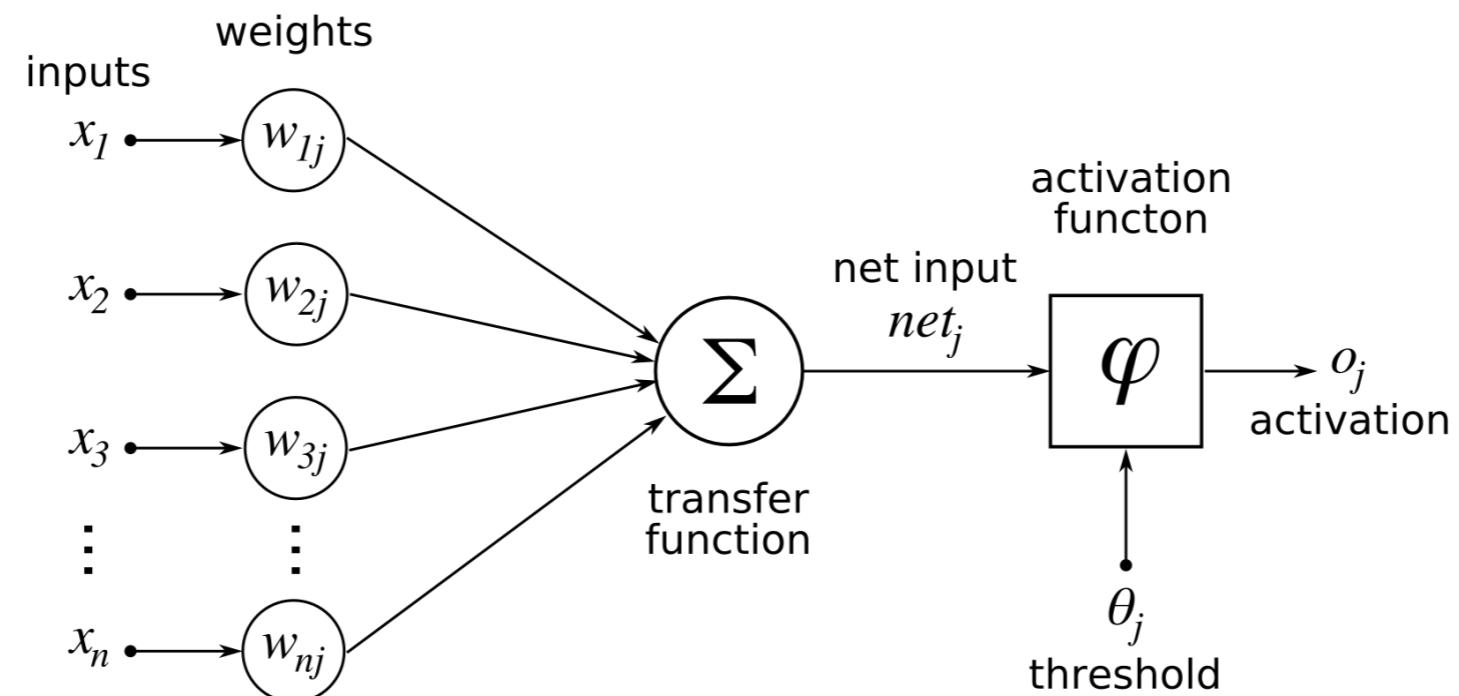
Neural Networks

- Feed-forward Network $y = \text{network}(x)$



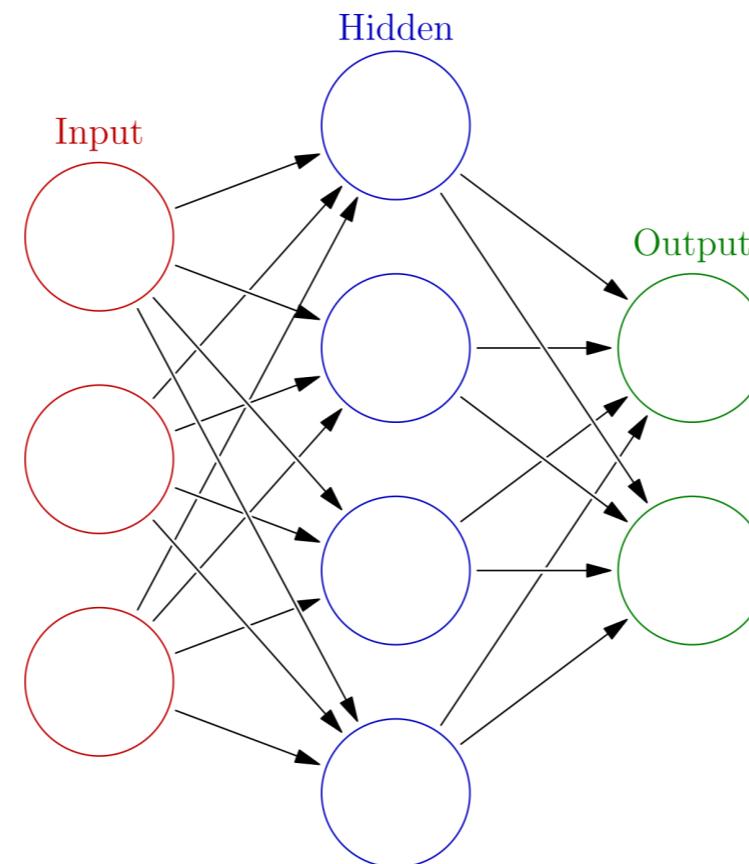
Neural Networks

- Feed-forward Network $y = \text{network}(x)$



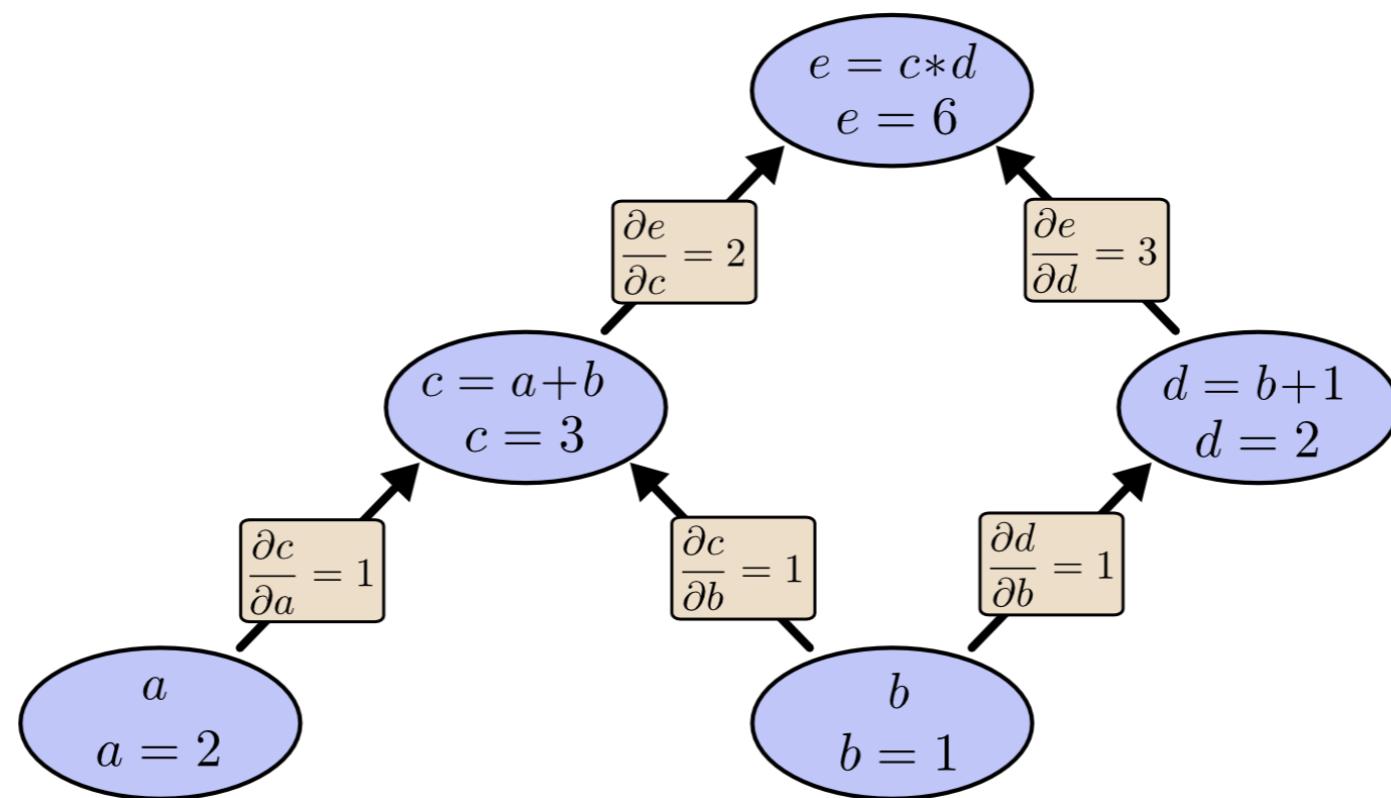
Neural Networks

- Feed-forward Network $y = \text{network}(x)$



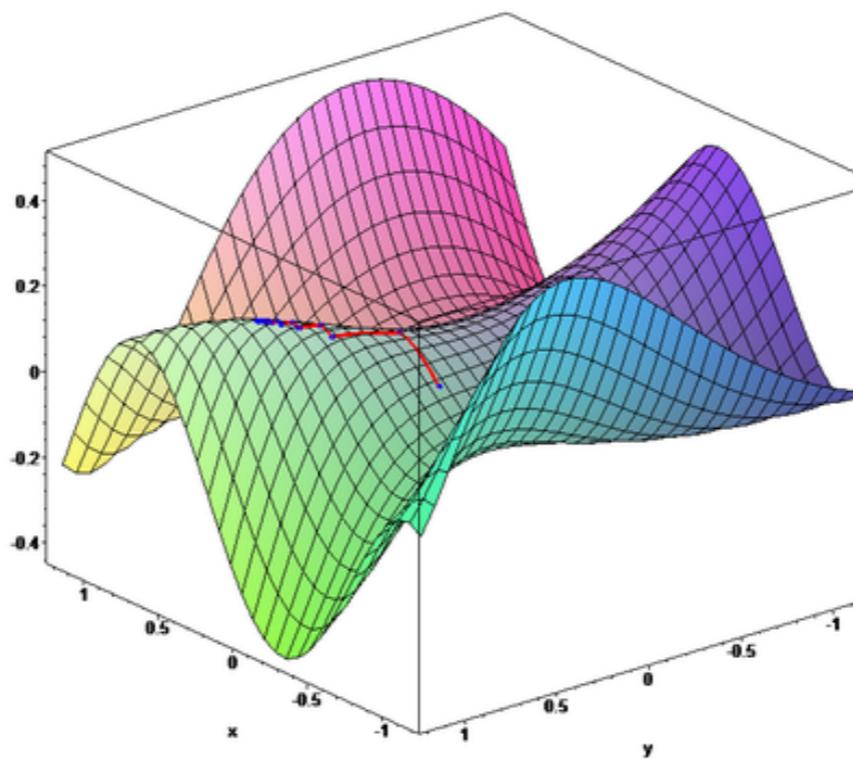
Neural Networks

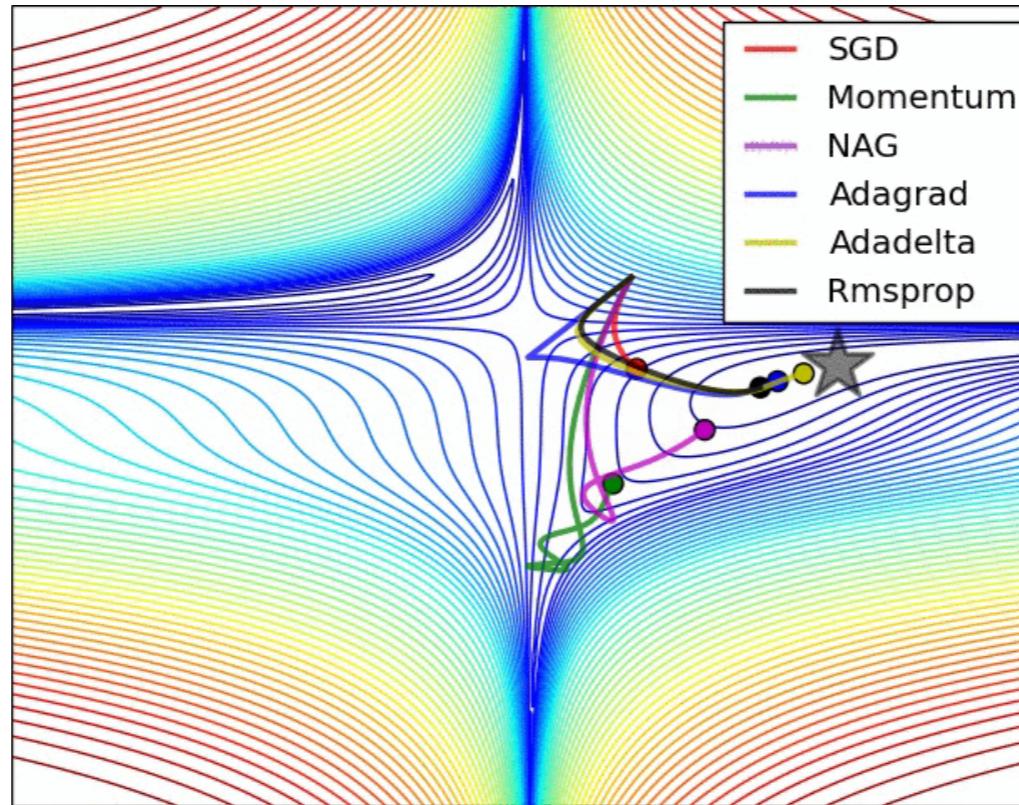
- Backpropagation `(loss_function(y_pred, y_true)).backward()`



Neural Networks

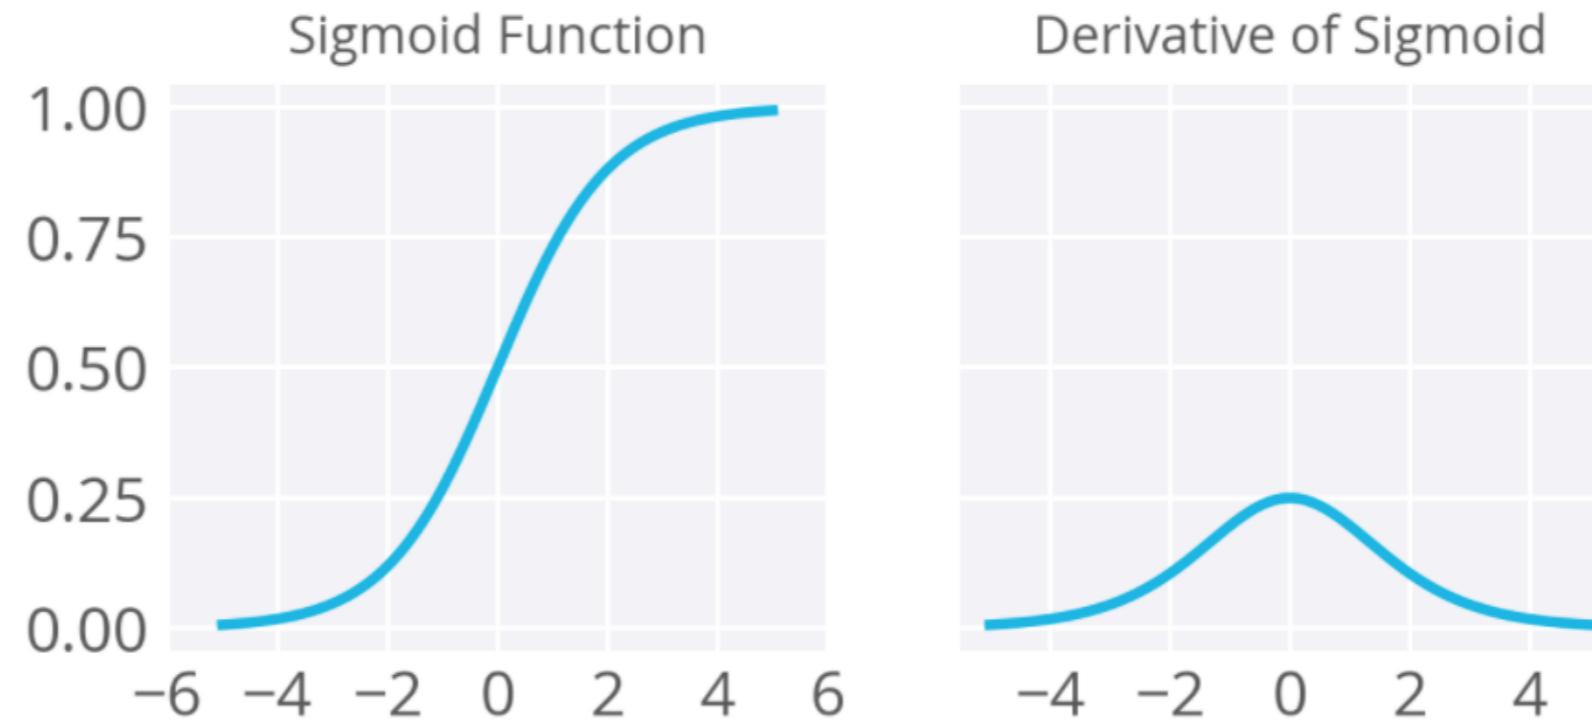
- Gradient Descent [torch/optim](#)
 - Stochastic Gradient Descent, Momentum, Adagrad, Adam, ...





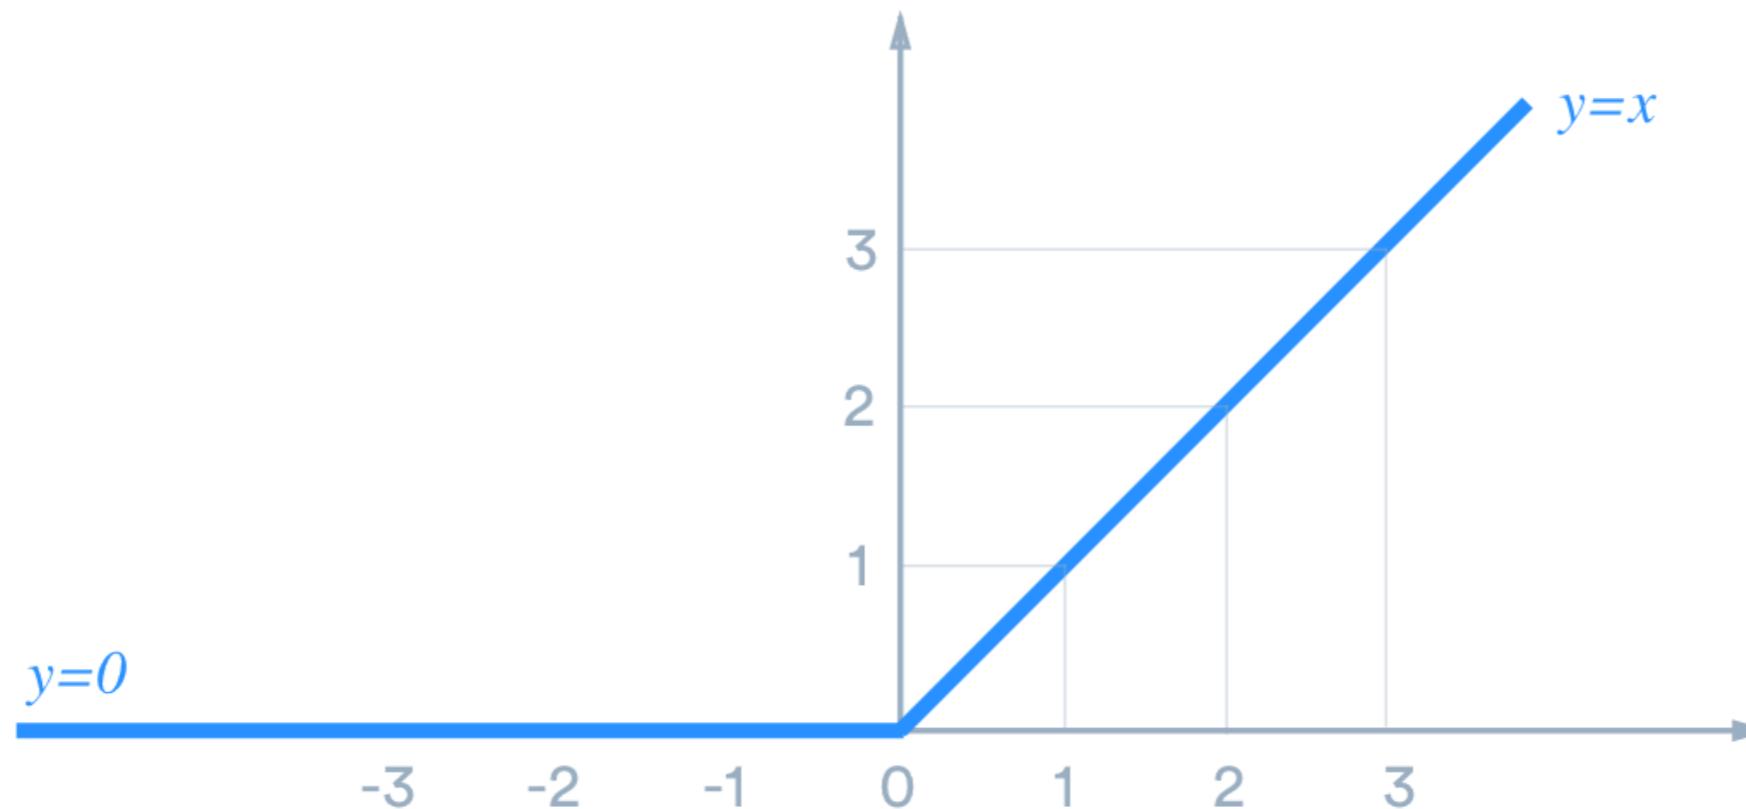
Activation functions and non-linearity

- $\text{sigmoid}(x) = \frac{1}{1+e^x}$ `torch.nn.Sigmoid`



Activation functions and non-linearity

- $relu(x) = \max(0, x)$ `torch.nn.ReLU`



Activation functions and non-linearity

- $\text{softmax}(z) = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}}$ `torch.nn.Softmax`



Loss functions

- L1 loss

$$L_1 = \sum_{i=1}^n |y_i - \hat{y}_i|$$

- L2 loss

$$L_2 = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Loss functions

- Mean Square Error `torch.nn.MSELoss`

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

- Cross Entropy `torch.nn.CrossEntropyLoss`

$$CE = - \sum_{i=1}^n y_i \ln(\hat{y}_i)$$

Loss functions

- Binary Cross Entropy `torch.nn.BCELoss`

$$BCE = -\frac{1}{n} \sum_{i=1}^n [y_i \ln(\hat{y}_i) + (1 - y_i) \ln(1 - \hat{y}_i)]$$

Regularization methods

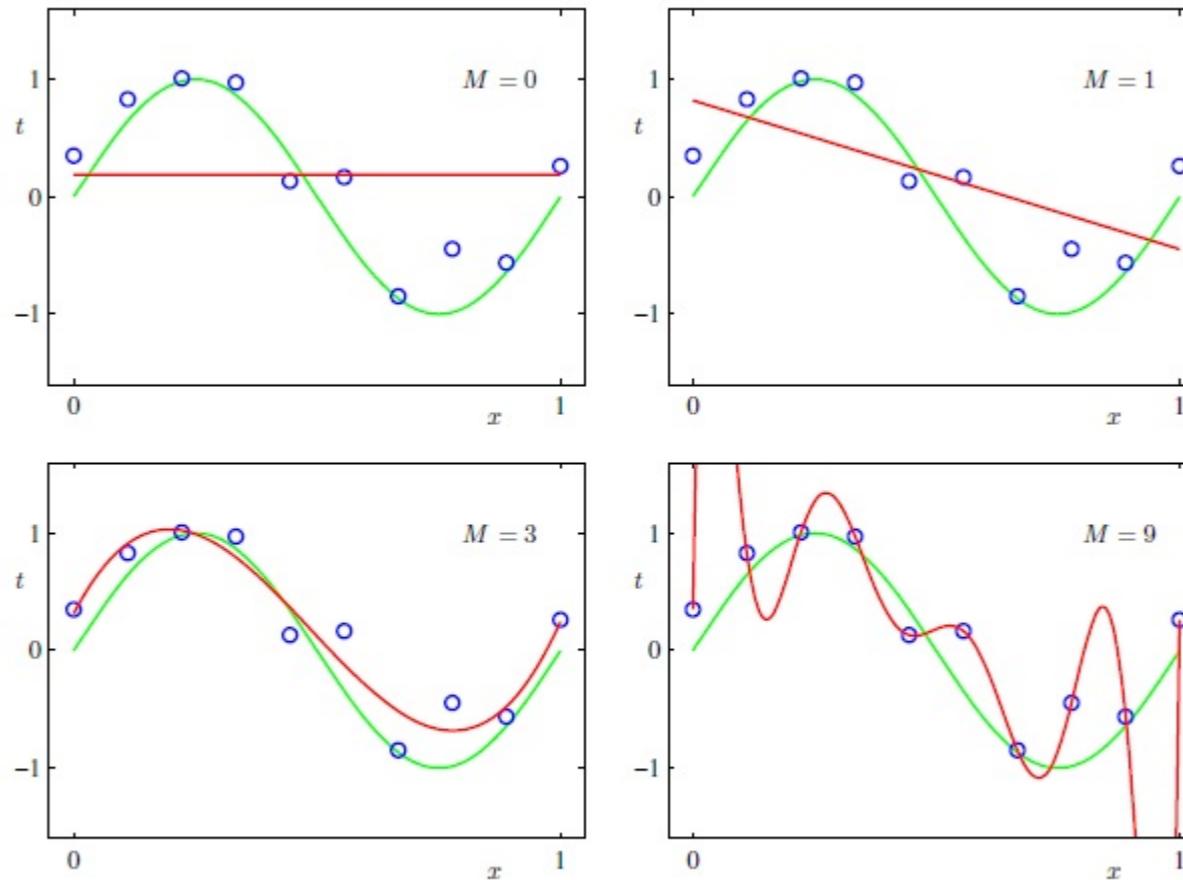


Figure 1.4 Plots of polynomials having various orders M , shown as red curves, fitted to the data set shown in Figure 1.2.

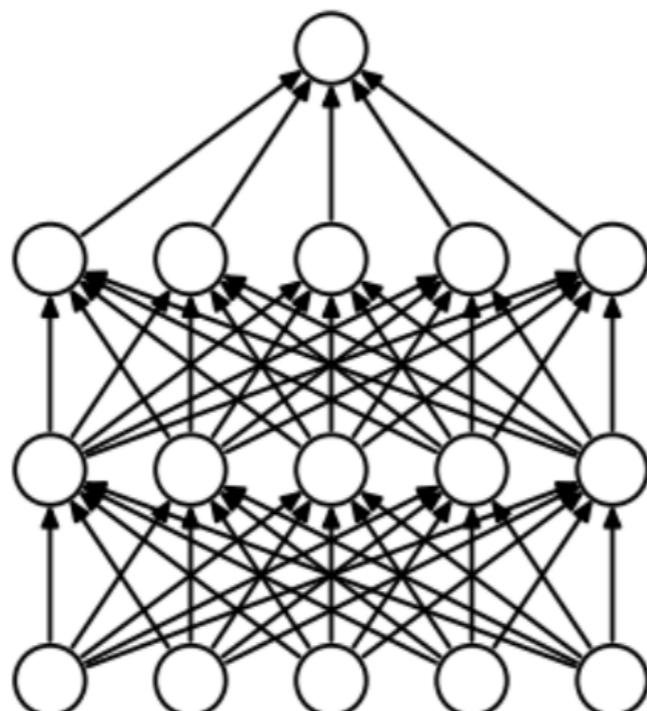
Regularization methods

- Weight decay

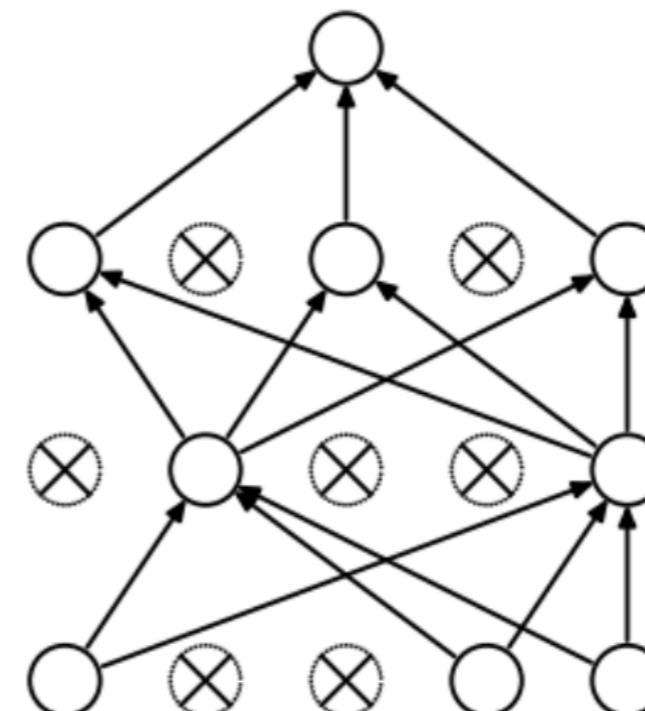
$$W \leftarrow W - \lambda \left(\frac{\partial L}{\partial W} + \gamma \|W\| \right)$$

Regularization methods

- Dropout



(a) Standard Neural Net



(b) After applying dropout.

2. Hello PyTorch





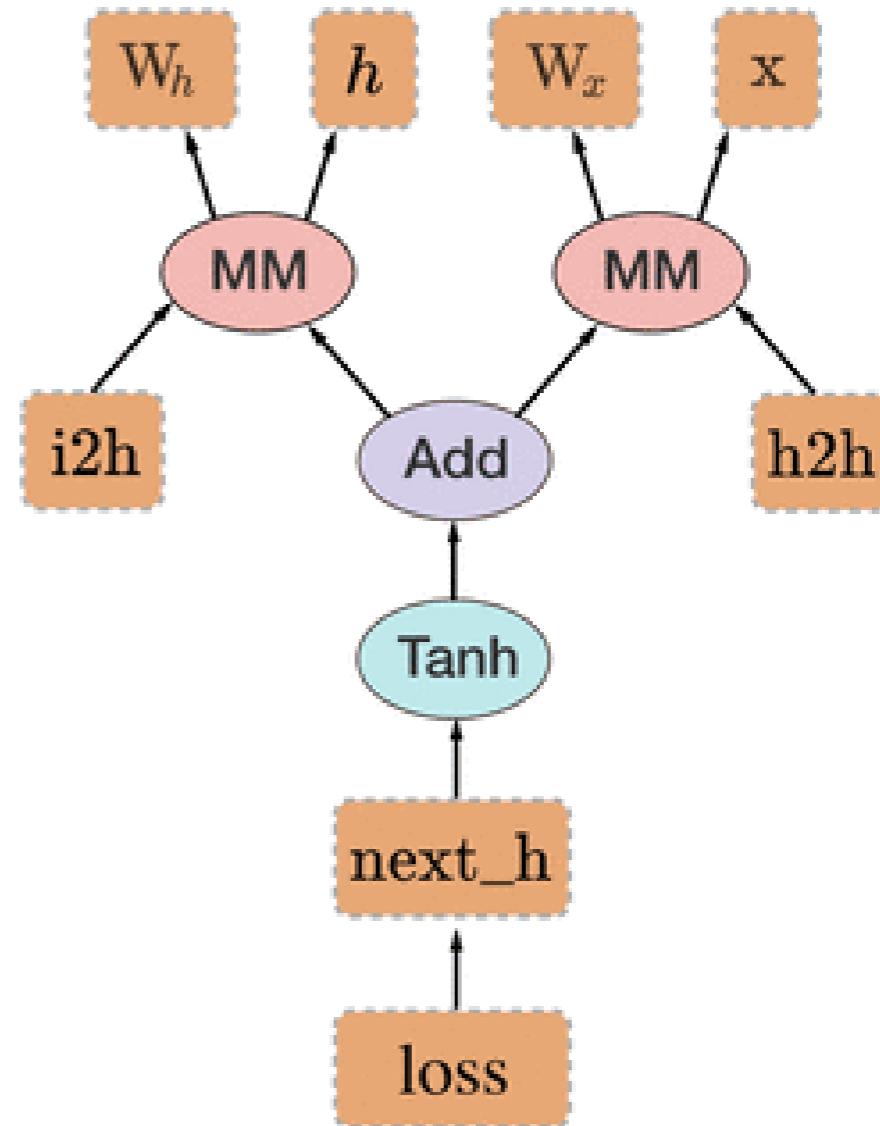
- Deep Learning Framework
 - Tensorflow, Keras, Torch, Chainer, MXNet
- Python-native, NumPy-friendly
- Dynamic graphs
- <https://pytorch.org/>
- <https://pytorch.org/docs/stable/index.html>

Back-propagation uses the dynamically created graph

```
x = torch.randn(1, 10)
prev_h = torch.randn(1, 20)
W_h = torch.randn(20, 20)
W_x = torch.randn(20, 10)

i2h = torch.mm(W_x, x.t())
h2h = torch.mm(W_h, prev_h.t())
next_h = i2h + h2h
next_h = next_h.tanh()

loss = next_h.sum()
loss.backward() # compute gradients!
```



Our stack

- **Ubuntu 16.04**

- Basic shell functions:

```
ls -lah , cd , cp , mv , cat , grep -irv , kill , ps -ef , history ,  
tar , unzip , apt , curl
```



Our stack

- **Python 3.6+**

The Zen of Python, by Tim Peters

Beautiful is better than ugly.

Explicit is better than **implicit**.

Simple is better than **complex**.

Complex is better than complicated.

Flat is better than nested.

Sparse is better than dense.

Readability counts.

Our stack

- **Conda**
 - Package manager + virtual environments
 - <https://conda.io/>



Our stack

- **Jupyter Notebook**
 - Document and visualize live code
 - <http://jupyter.org/>



Our stack

- **git**
 - Version control system
 - Basic commands:
`git clone` , `git checkout` , `git add` , `git commit` , `git push` , `git pull` , `git diff` , `git log`
 - <https://github.com/> , <https://gitlab.com/>

Our stack

- Docker CE
 - Container virtualization
 - Basic commands:
`docker run , docker exec , docker start , docker stop`

